


# 30 SQL Queries Interview Questions and Answers



Writing SQL queries during interviews can be tough for beginners and experienced developers. However, with the right preparation and practice, you can ace your next SQL queries interview.

In this guide, I'll walk you through both basic and advanced SQL queries like SELECT, JOIN, GROUP BY, and window functions. Each answer will be short and clear and include SQL code examples for you to understand the concepts better. By the end of this guide, you'll be in a much better position to ace your SQL queries interview and build solid skills you can use beyond it.



Preparing for your SQL queries interview

Test yourself with Flashcards

Questions List

+

I have included a set of flashcards to help you study and practice more efficiently. If you are just starting out in your career, checkout roadmap.sh's [SQL roadmap](#) or take our [SQL queries course](#) to master this topic in under 60 hours!

## Preparing for your SQL queries interview

While preparing for your interview, you should remember the following points.

- Review the basics of SQL. You should know what SQL stands for and what it is used for.
- Make sure you have a basic understanding of databases and the different types of databases, such as SQL and NoSQL databases.
- Consider reading about the SQL data types and basic database concepts such as indexing, foreign keys, primary keys, etc.
- Practice writing SQL queries on local databases or online platforms like [HackerRank](#).
- Get familiar with at least one relational database management system such as [PostgreSQL](#), Microsoft SQL Server, MySQL, Oracle DB.
- On a general note, read up on the company you are interviewing with to know more about what they do, and also prepare some questions

beforehand to show you are interested in what they do.

## Test yourself with Flashcards

You can either use these flashcards or jump to the questions list section below to see them in a list format.

< 0 / 30 >

✓ Knew 0 Items ✨ Learnt 0 Items ▶ Skipped 0 Items 🔁 **ResetProgress**

### Foundational SQL Queries

**What is the difference between  
WHERE and HAVING?**

[Click to Reveal the Answer](#)

✓ Already Know that

✨ Didn't Know that

▶ **Skip Question**

## Questions List

If you prefer to see the questions in a list format, you can find them below.

### Foundational SQL Queries

**What is the difference between WHERE and HAVING?**

You use **WHERE** for filtering rows before applying any grouping or aggregation. The code snippet below illustrates the use of **WHERE**. It filters the `Users` table for rows where the `Age` is greater than 18.

sql



```
SELECT * FROM Users
WHERE Age > 18;
```

The result of the query is similar to the table below.

| userId | firstName | lastName | age |
|--------|-----------|----------|-----|
| 1      | John      | Doe      | 30  |
| 2      | Jane      | Don      | 31  |
| 3      | Will      | Liam     | 25  |
| 4      | Wade      | Great    | 32  |
| 5      | Peter     | Smith    | 27  |

On the other hand, you use **HAVING** to filter groups after performing grouping and aggregation. You apply it to the result of aggregate functions, and it is mostly used with the **GROUP BY** clause.

sql



```
SELECT FirstName, Age FROM Users
GROUP BY FirstName, Age
HAVING Age > 30;
```

The code above selects the `FirstName` and `Age` columns, then groups by the `FirstName` and `Age`, and finally gets entries with age greater than 30. The result of the query looks like this:

| firstName | age |
|-----------|-----|
| Wade      | 32  |
| Jane      | 31  |

## How do you find duplicates in a table?

To find duplicate records, you must first define the criteria for detecting duplicates. Is it a combination of two or more columns where you want to detect the duplicates, or are you searching for duplicates within a single column?

The following steps will help you find duplicate data in a table.

- Use the **GROUP BY** clause to group all the rows by the column(s) on which you want to check the duplicate values.
- Use the **COUNT** function in the **HAVING** command to check if any groups have more than one entry.

Let's see how to handle single-column duplicates. In a table `Users`, there are three users who are 30 years of age. Let's use the **GROUP BY** clause and **COUNT** function to find the duplicate values.

sql



```
SELECT Age, COUNT(Age)
FROM Users
```

```
GROUP BY Age
HAVING COUNT(Age) > 1
```

The result of the query looks like this:

| age | count |
|-----|-------|
| 30  | 3     |

Handling multi-column (composite) duplicates is similar to handling single-column duplicates.

sql



```
SELECT FirstName, LastName, COUNT(*) AS dup_count
FROM Users
GROUP BY FirstName, LastName
HAVING COUNT(*) > 1;
```

After finding duplicates, you might be asked how to delete the duplicates. The query to delete duplicates is shown below using Common Table Expression (CTE) and ROW\_NUMBER().

sql



```
WITH ranked AS (  
    SELECT *,  
           ROW_NUMBER() OVER (PARTITION BY Age ORDER BY id) AS rn  
    FROM Users  
)  
DELETE FROM Users  
WHERE id IN (  
    SELECT id  
    FROM ranked  
    WHERE rn > 1  
);
```

The query deletes all the duplicates while retaining the first row of data.


## What is the difference between INNER JOIN and LEFT JOIN?

A **JOIN** combines data from two or more tables based on a related column between them. It is useful when you need to retrieve data spread across multiple tables in relational database management systems.


An **INNER JOIN** returns only rows with a match in both tables based on the specified join condition. If there are no matching rows, there will be no results. The SQL syntax for an **INNER JOIN** is shown in the code snippet below.

# Inner join vs. Left join

Inner join



Left join



sql



```
SELECT table1.column_name1, table1.column_name2, table2.column_name1, table2.column_n  
INNER JOIN table2  
ON table1.column_name = table2.column_name
```

For example, there are two tables `Users` and `Cities` with the following data:

Users table



| userId | firstName | lastName | age | cityId |
|--------|-----------|----------|-----|--------|
| 1      | John      | Doe      | 30  | 1      |
| 2      | Jane      | Don      | 31  | 1      |
| 3      | Will      | Liam     | 25  | 1      |
| 4      | Wade      | Great    | 32  | 1      |
| 5      | Peter     | Smith    | 27  | 2      |
| 6      | Rich      | Mond     | 30  | 2      |
| 7      | Rach      | Mane     | 30  | 2      |
| 8      | Zach      | Ridge    | 30  | 3      |

## Cities table

| id | name       |
|----|------------|
| 1  | London     |
| 2  | Manchester |

Let's say you want to retrieve a list of users and their respective city names. You can achieve this using the **INNER JOIN** query.

sql



```
SELECT users.firstName, users.lastName, users.age, cities.name as cityName FROM users
INNER JOIN cities
ON users.cityId = cities.id
```

| firstName | lastName | age | cityName   |
|-----------|----------|-----|------------|
| John      | Doe      | 30  | London     |
| Jane      | Don      | 31  | London     |
| Will      | Liam     | 25  | London     |
| Wade      | Great    | 32  | London     |
| Peter     | Smith    | 27  | Manchester |
| Rich      | Mond     | 30  | Manchester |
| Rach      | Mane     | 30  | Manchester |

**LEFT JOIN** returns all the rows from the left table (table 1) and the matched rows from the right table (table 2). If no matching rows exist in the right table (table 2), then NULL values are returned. The SQL syntax for a Left join is shown in the code snippet below.

sql



```
SELECT table1.column_name1, table1.column_name2, table2.column_name1, table2.column_n  
LEFT JOIN table2  
ON table1.column_name = table2.column_name
```

Let's have a look at a practical example with `Users`` and `Cities`` tables from before.

When you execute the **LEFT JOIN** query, you get the table below.

| firstName | lastName | age | cityName   |
|-----------|----------|-----|------------|
| John      | Doe      | 30  | London     |
| Jane      | Don      | 31  | London     |
| Will      | Liam     | 25  | London     |
| Wade      | Great    | 32  | London     |
| Peter     | Smith    | 27  | Manchester |
| Rich      | Mond     | 30  | Manchester |
| Rach      | Mane     | 30  | Manchester |
| Zach      | Ridge    | 30  | null       |

Write a query to find the second highest salary from a table

Given a table ``Salaries``,

| id | salary |
|----|--------|
| 1  | 1000   |
| 2  | 2000   |
| 3  | 3000   |
| 4  | 4000   |

The query to find the second-highest salary is shown in the code snippet below

sql

```
SELECT DISTINCT Salary
FROM Salaries
ORDER BY Salary DESC
LIMIT 1 OFFSET 1
```

The result of the query is shown below


|   | salary |
|---|--------|
| 1 | 3000   |

## What is the difference between UNION and UNION ALL?

UNION is used for removing duplicates while UNION ALL keeps all duplicates. UNION is slower compared to UNION ALL because of de-duplication. You use UNION when you want to obtain unique records and UNION ALL when you want every row even if they are repeated.

## What are indexes and why are they useful?

Indexes in databases are like the indexes in books. They increase the speed of data retrieval from a database. When you want to read data from a table, instead of going through all the rows of the table, indexes help to go straight to the row you are looking for.



They improve **SELECT** queries, improve performance, and make sorting and filtering faster. They also ensure data integrity. There are different types of indexes, which include:

- B-Tree index
- Composite index
- Unique index
- Full text index

- Bitmap index
- Clustered index
- Non-clustered index

## What is a primary key?

A primary key is the unique identifier of a row of data in a table. You use it to identify each row uniquely, and no two rows can have the same primary key. A primary key column cannot be null. In the example below, `user_id` is the primary key.

sql



```
CREATE TABLE users (  
  user_id INT PRIMARY KEY,  
  name VARCHAR(100),  
  phoneNumber VARCHAR(100)  
);
```

## What is a foreign key?

A foreign key is like a bridge between two tables. A foreign key in one table is the primary key in another. It is the connector between the two tables.

## Aggregation And Grouping

### How does GROUP BY work?

**GROUP BY** is a standard SQL command that groups rows with the same value in the specified column. You should use with aggregate functions such as **COUNT**, **MIN**, **MAX**, etc.

## Group by



The query below illustrates the **GROUP BY** clause:

sql



```
SELECT columnName FROM Table
GROUP BY columnName
```

What happens if you SELECT a column not in the GROUP BY clause?

If you **SELECT** a column not in the **GROUP BY** clause, it will throw an error stating that the column must be in the **GROUP BY** clause or in an aggregate function. Let's use the table below as an illustration.

| firstName | lastName  | phoneNumber |
|-----------|-----------|-------------|
| John      | Doe       | +23410910   |
| Jack      | Ray       | +23410911   |
| Irene     | Rotherdam | +23410911   |

If you run the query below against the database:

sql

```
SELECT firstName, phoneNumber FROM phoneNumbers
GROUP BY phoneNumber
```

The result will be an error because `firstName` is not in the **GROUP BY** clause and not using an aggregate function.

**Write a query to COUNT the number of users by country**

Given a table `Users` that looks like this:

| userId | firstName | lastName | age | country  |
|--------|-----------|----------|-----|----------|
| 1      | John      | Doe      | 30  | Portugal |
| 2      | Jane      | Don      | 31  | Belgium  |



| userId | firstName | lastName | age | country   |
|--------|-----------|----------|-----|-----------|
| 3      | Will      | Liam     | 25  | Argentina |
| 4      | Wade      | Great    | 32  | Denmark   |
| 5      | Peter     | Smith    | 27  | USA       |
| 6      | Rich      | Mond     | 30  | USA       |
| 7      | Rach      | Mane     | 30  | Argentina |
| 8      | Zach      | Ridge    | 30  | Portugal  |

The query to **COUNT** the number of users by country is:

sql

```
SELECT country, COUNT(country) FROM users
GROUP BY country
```

The query uses the **GROUP BY** clause to group the users by country and then shows the count in the next column. The result of the query looks like this:

| country   | count |
|-----------|-------|
| USA       | 2     |
| Portugal  | 2     |
| Argentina | 2     |

| country | count |
|---------|-------|
| Belgium | 1     |
| Denmark | 1     |

## What happens if you use GROUP BY without an aggregate function?

If you use the **GROUP BY** clause without an aggregate function, it is equivalent to using the **DISTINCT** command. For example, the command below:

sql



```
SELECT phoneNumber FROM phoneNumbers  
GROUP BY phoneNumber
```

is equivalent to:

sql



```
SELECT DISTINCT phoneNumber FROM phoneNumbers
```

## What is the difference between COUNT(\*) and COUNT(column\_name)?

The difference is that **COUNT(\*)** counts all the rows of data, including NULL values, while **COUNT(column\_name)** counts only non-NULL values in the specified column. Let's illustrate this using a table named `Users``.

| userId | firstName | lastName | age | country  |
|--------|-----------|----------|-----|----------|
| 1      | John      | Doe      | 30  | Portugal |
| 2      | Jane      | Don      | 31  | Belgium  |
| 3      | Zach      | Ridge    | 30  | Norway   |
| 4      | null      | Tom      | 25  | Denmark  |


If you use **COUNT(\*)**, the result will be 4 but if you use **COUNT(firstName)**, it will return 3, omitting the null value.

## Subqueries And Nested Logic

### What is the difference between a subquery and a JOIN?

A subquery is a query that is inside another query. You use it for queries that require complex logic. You should use subqueries when you want to use the result of that subquery for another query. In the example below, the subquery is in brackets.

# SQL subquery



 roadmap.sh

sql



```
SELECT firstName,  
  (SELECT COUNT(*)  
   FROM cities  
   WHERE cities.id = users.city_id) AS cityCount  
FROM users;
```

On the other hand, a **JOIN** combines two or more tables based on related columns between them. The related column is usually a foreign key. You

should use **JOINS** when you want to pull related data from different tables together. The code below illustrates how to use a **JOIN**.

sql



```
SELECT firstName, COUNT(*) FROM users  
JOIN cities ON users.city_id = cities.id
```

A JOIN is faster than a subquery in the following scenarios:

- When you are querying data from multiple tables.
- When you are filtering or joining on index columns.

**Write a query to find employees earning more than the average salary**

Given an `Employees` table with columns `id`, `name`, and `salary` that looks like this:

| id | name    | salary |
|----|---------|--------|
| 1  | Irene   | 1000   |
| 2  | Peter   | 1230   |
| 3  | Raymond | 1450   |
| 4  | Henry   | 1790   |
| 5  | Naomi   | 2350   |
| 6  | Bridget | 2000   |

| id | name     | salary |
|----|----------|--------|
| 7  | Emily    | 2500   |
| 8  | Great    | 3000   |
| 9  | Mercedes | 2750   |
| 10 | Zoe      | 2900   |

The query to find employees earning more than the average salary is:

sql



```
SELECT * FROM employees
WHERE salary > (SELECT AVG(salary) FROM employees);
```

| id | name     | salary |
|----|----------|--------|
| 5  | Naomi    | 2350   |
| 7  | Emily    | 2500   |
| 8  | Great    | 3000   |
| 9  | Mercedes | 2750   |
| 10 | Zoe      | 2900   |

## Explain how a correlated subquery works

A correlated subquery is a subquery that depends on a value from the outer query. This means that the query is evaluated for each row that might

be selected in the outer query. Below is an example of a correlated subquery.

sql



```
SELECT name, country_id, salary
FROM employees em
WHERE salary > (
    SELECT AVG(salary) FROM employees
    country_id = em.country_id);
```

The code above:

- Runs the outer query through each row of the table.
- Takes the `country_id` from the `employees` table.
- Iterates through the other rows and does the same calculation.

This leads to a degrading performance as the data in the table grows.

You should use a correlated subquery if you want to perform row-specific operations or cannot achieve an operation using JOIN or other aggregate functions.

### When should you use EXISTS instead of IN in a subquery?

**EXISTS** and **IN** are used in subqueries to filter results, but they perform different functions depending on their usage.

You should use **EXISTS** in the following situations:

- When you want to check if a row exists and not the actual values.
- When the subquery is a correlated query.

- When the subquery returns many rows but you want to get the first match.

You should use **IN** in the following scenarios:

- When you are comparing a column to a list of values.
- When the subquery returns a small or static list.

### Can you nest subqueries multiple levels deep?

Yes, you can nest subqueries multiple levels deep when you want to perform complex logic. A nested subquery is a subquery inside another subquery, forming layers of subqueries. Many SQL engines allow multiple layers of subqueries, but this causes poor readability and degrades performance.

## Window Functions And Advanced Queries

### What is a window function?

A window function is a function that allows you to perform operations on a specific set of rows related to the current row. Unlike aggregate functions that perform calculations on an entire data set, window functions can perform operations on a subset of data. These calculations are valid for aggregates, ranking, and cumulative totals without altering the original dataset.

### Write a query to calculate a running total

Let's use a table ``Sales`` as a reference for this query. It has three columns: ``id``, ``day`` which represents the day of the week, and ``amount`` which is



the amount sold in US Dollars. The table looks like this:

| id | day       | amount |
|----|-----------|--------|
| 1  | Monday    | 200    |
| 2  | Tuesday   | 300    |
| 3  | Wednesday | 600    |
| 4  | Thursday  | 390    |
| 5  | Friday    | 900    |

The query to calculate the running total is:

sql



```
SELECT
  id,
  sale_date,
  amount,
  SUM(amount) OVER (ORDER BY sale_date) AS running_total
FROM
  sales;
```

The query uses a Window function **OVER** to sum the amount for each row of data and saving the running total. It gets the total for each day and adds it to the previous totals. The result of the query looks like this:

| id | day    | amount | running_total |
|----|--------|--------|---------------|
| 1  | Monday | 200    | 200           |

| id | day       | amount | running_total |
|----|-----------|--------|---------------|
| 2  | Tuesday   | 300    | 500           |
| 4  | Thursday  | 390    | 1100          |
| 3  | Wednesday | 600    | 1490          |
| 5  | Friday    | 900    | 2390          |

You can observe from the image that the last column is `running_total`, which takes the amount for the current day and adds it to its previous value to get its current value.

**What is the difference between `RANK()`, `DENSE_RANK()`, and `ROW_NUMBER()`?**

The **`RANK()`** function assigns each row a rank according to an ascending or descending order. If there are matching values, it assigns them the same position and then skips the next number for the next rank. For example, if two rows have equivalent values and are both assigned rank 1, the next rank would be 3 instead of 2.

# Window functions



 roadmap.sh

Let's use the `Sales` table from the previous question to illustrate the `RANK()` function. The query to rank in order of the amount looks like this:

sql



```
SELECT
  id,
  day,
  amount,
  RANK() OVER (ORDER BY amount DESC) AS amount_rank
FROM
  sales;
```

The result is shown in the image below. You will observe that the amount 900 takes the first rank and 200 the lowest rank. Also, there is a gap between rank 2 and 4 because two values have the same rank. You can also infer that the most sales were on Friday and the least on Monday.

| id | day       | amount | amount_rank |
|----|-----------|--------|-------------|
| 5  | Friday    | 900    | 1           |
| 3  | Wednesday | 600    | 2           |
| 6  | Saturday  | 600    | 2           |
| 4  | Thursday  | 390    | 4           |
| 2  | Tuesday   | 300    | 5           |
| 1  | Monday    | 200    | 6           |

**DENSE\_RANK()** function is similar to **RANK()** in that it assigns ranks to rows, but the difference is that **DENSE\_RANK** does not leave a gap when there are two or more equivalent values. Let's illustrate it with the ``Sales`` table from above. The query is shown below.

sql



```
SELECT
  id,
  day,
  amount,
  DENSE_RANK() OVER (ORDER BY amount DESC) AS amount_rank
FROM
  sales;
```

The result is shown below. As you will notice, there is no gap between the ranks like in the **RANK** function.

| id | day       | amount | amount_rank |
|----|-----------|--------|-------------|
| 5  | Friday    | 900    | 1           |
| 3  | Wednesday | 600    | 2           |
| 6  | Saturday  | 600    | 2           |
| 4  | Thursday  | 390    | 3           |
| 2  | Tuesday   | 300    | 4           |
| 1  | Monday    | 200    | 5           |

**ROW\_NUMBER** assigns a unique number to each row depending on the order you specify. It does not skip numbers; even though there are equivalent values, it assigns them different numbers, unlike **RANK** and **DENSE\_RANK** functions that give them the same rank.

Let's use the same ``Sales`` table to illustrate. The query below shows how to use the **ROW\_NUMBER** function.

sql



```
SELECT
  id,
  day,
  amount,
  ROW_NUMBER() OVER (ORDER BY amount DESC) AS rowNumber
FROM
  sales;
```

The result is shown in the image below. You will notice that the `rownumber`` column increases, and even though there are matching values, it just assigns a unique row number to each.

| id | day       | amount | amount_rank |
|----|-----------|--------|-------------|
| 5  | Friday    | 900    | 1           |
| 3  | Wednesday | 600    | 2           |
| 6  | Saturday  | 600    | 3           |
| 4  | Thursday  | 390    | 4           |
| 2  | Tuesday   | 300    | 5           |
| 1  | Monday    | 200    | 6           |

## What is LAG() and LEAD() in SQL? Give an example use case

**LAG()** and **LEAD()** are window functions used to retrieve data from rows before and after a specified row. You can also refer to them as positional SQL functions.

**LAG()** allows you to access a value stored in rows before the current row. The row may be directly before or some rows before. Let's take a look at the syntax:

sql



```
LAG(column_name, offset, default_value)
```

It takes three arguments.

- **column\_name**: This specifies the column to fetch from the previous row.
- **offset**: This is an optional argument and specifies the number of rows behind to look at. The default is 1.
- **default\_value**: This is the value to assign when no previous row exists. It is optional, and the default is NULL.

Using the ``Sales`` table, let's illustrate the **LAG()** function. The query is used to find the previous day sales. LAG() is useful when you want to create reports of past events.

| id | day       | amount |
|----|-----------|--------|
| 1  | Monday    | 200    |
| 2  | Tuesday   | 300    |
| 3  | Wednesday | 600    |
| 4  | Thursday  | 390    |
| 5  | Friday    | 900    |
| 6  | Saturday  | 600    |

sql



```
SELECT
  id,
  day,
  amount,
```

```
LAG(amount) OVER (ORDER BY id) AS previous_day_sales
FROM
sales;
```

The result of the query looks like this:

| id | day       | amount | previous_day_sales |
|----|-----------|--------|--------------------|
| 1  | Monday    | 200    | null               |
| 2  | Tuesday   | 300    | 200                |
| 3  | Wednesday | 600    | 300                |
| 4  | Thursday  | 390    | 600                |
| 5  | Friday    | 900    | 390                |
| 6  | Saturday  | 600    | 900                |

You use the **LEAD()** function to get data from rows after the current row. Its syntax is similar to that of the **LAG()** function. You can use it for forecasting future trends by looking ahead.

The query using the **LEAD()** function is shown below.

sql





```
SELECT
  id,
  day,
  amount,
  LEAD(amount) OVER (ORDER BY id) AS previous_day_sales
FROM
  sales;
```

| id | day       | amount | previous_day_sales |
|----|-----------|--------|--------------------|
| 1  | Monday    | 200    | 300                |
| 2  | Tuesday   | 300    | 600                |
| 3  | Wednesday | 600    | 390                |
| 4  | Thursday  | 390    | 900                |
| 5  | Friday    | 900    | 600                |
| 6  | Saturday  | 600    | null               |

## How will you detect gaps in a sequence of dates per user?

You will use the **LAG()** function to detect gaps in a sequence of dates per user. You will compare each date with the previous one and check if the difference is greater than 1.

Let's use a table ``ClockIns`` to demonstrate how you detect gaps. The table has two columns, ``userId`` and ``clockInDate``, representing the user identification number and the date the user clocked in with an access card into a facility. The table looks like this:

| userId | clockInDate |
|--------|-------------|
| 1      | 2025-01-01  |
| 1      | 2025-01-02  |
| 1      | 2025-01-05  |
| 1      | 2025-01-06  |
| 2      | 2025-01-06  |
| 2      | 2025-01-06  |
| 2      | 2025-01-07  |
| 3      | 2025-01-02  |
| 3      | 2025-01-04  |
| 3      | 2025-01-06  |
| 3      | 2025-01-07  |

To query to find gaps per user looks like this:

sql



```
WITH clockInGaps AS (  
  SELECT  
    userid,  
    clockInDate,  
    LAG(clockInDate) OVER (PARTITION BY userId ORDER BY clockInDate) AS previousClock  
  FROM  
    clockIns  
)
```

```
SELECT
  userId,
  previousClockInDate AS gapStart,
  clockInDate AS gapEnd,
  clockInDate - previousClockInDate - 1 AS gapDays
FROM clockInGaps
WHERE clockInDate - previousClockInDate > 1
ORDER BY userId, gapStart;
```

The code above starts with creating an expression `clockInGaps` that queries for each user and their `clockInDate` and uses the `LAG` function to get the previous date for each user. Then, the main query filters each row and finds the gaps between the current date and the previous date. The result of the query looks like this:

| userId | gapStart   | gapEnd     | gapDays |
|--------|------------|------------|---------|
| 1      | 2025-01-02 | 2025-01-05 | 2       |
| 2      | 2025-01-07 | 2025-01-10 | 2       |
| 3      | 2025-01-02 | 2025-01-04 | 1       |
| 3      | 2025-01-04 | 2025-01-06 | 1       |

**What does the `NTILE()` function do, and how might it be useful in analyzing data?**

**NTILE()** is a window function that divides rows into a pre-defined number of roughly equal groups. It's like breaking your data into different sets based on your defined criteria. For example, let's say you have some student scores from 1 to 100; you can use the **NTILE()** function to categorize the scores into different groups or buckets.

The syntax of the `NTILE()` function is:

sql



```
NTILE(n) OVER (ORDER BY some_column)
```

- `n`: represents the number of groups you want to divide your rows into.
- `ORDER BY`: defines the order of the rows in each group where the function is applied.

Let's see a practical example using a table `Scores`. The table stores students' scores on a test. We will see how to use the `NTILE()` function.

| userId | score |
|--------|-------|
| 1      | 78    |
| 2      | 70    |
| 3      | 90    |
| 4      | 98    |
| 5      | 60    |
| 6      | 88    |
| 7      | 100   |
| 8      | 66    |

The query using the `NTILE()` function looks like this:

sql



```
SELECT
  id,
  score,
  NTILE(3) OVER (ORDER BY score DESC) AS category
FROM scores;
```

| userId | score | category |
|--------|-------|----------|
| 7      | 100   | 1        |
| 4      | 98    | 1        |
| 3      | 90    | 1        |
| 6      | 88    | 2        |
| 1      | 78    | 2        |
| 2      | 70    | 2        |
| 8      | 66    | 3        |
| 5      | 60    | 3        |

The **NTILE()** function is useful in data analysis because it can detect outliers in a data set and create histograms of data. It can also create percentiles and quartiles for data distribution.

## Optimization And Pitfalls

How would you optimize slow-running queries?

To optimize slow-running queries, you need to analyze the query first to know what to optimize. You can perform different optimizations depending on the query. Some of the optimizations include:

- **Using indexes effectively:** Indexes speed up queries by enabling the database to find entries that fit specific criteria quickly. Indexing is the process of mapping the values of one or more columns to a unique value that makes it easy to search for rows that match a search criteria. You can create indexes on columns used frequently in the **WHERE**, **JOIN**, and **ORDER BY** clauses. However, note that creating too many indexes can slow down inserts, updates, and deletions.
- **\*\*Avoid SELECT \*\***: Using the **SELECT \*\*\*\*\*** statement can slow down your query performance because it returns all the columns in a table including the ones not needed for the query. You should select only the columns that you need for a query for optimal performance. So when you see a query that selects all columns, you should check if all the columns are really needed and used further down the query chain.
- **Avoid using subqueries:** Subqueries slow down query performance, especially when you use them in the **WHERE** or **HAVING** clauses. You should avoid using subqueries where possible and use JOINS or other techniques instead.
- **Utilize stored procedures:** Stored procedures are precompiled SQL statements stored in a database, and can be called from an application or directly from a query. Using stored procedures can improve your query performance by reducing the amount of data that is sent between the database and your application, and also saves time required to compile the SQL statements.

## Why should you avoid SELECT \* in production code?

You should avoid using `**SELECT **` as much as possible in your production code for the following reasons:

- **Increased IO:** Using `**SELECT **`, you can return unnecessary data that leads to increased Input/Output cycles at the database level since you will be reading all the data in a table. This effect will be more impactful on a table with a lot of data and even slow down your query.
- **Increased network traffic:** `**SELECT **` returns more data than required to the client, which uses more network bandwidth than needed. The increase in network bandwidth causes data to take longer to reach the client application and impacts the application's performance.
- **More application memory:** The return of a lot of data would make your application require more memory to hold the unnecessary data which might you might not use and this impacts application performance.
- **Makes maintenance more difficult:** Using `**SELECT **` makes code maintenance more challenging. If the table structure changes by adding, removing, or renaming columns, the queries using `**SELECT **` could break unexpectedly. You should explicitly specify the columns from which you want to fetch data to ensure resilience against potential changes in the database schema.

## What is the impact of missing indexes?

Missing indexes can affect the performance of queries, especially when the data grows. The major impacts of missing indexes are listed below:

- **Slow queries:** Without indexes, every read query will go through the whole table to find matching rows. This will get worse as the data in the table grows.
- **Locking and concurrency issues:** Scanning a table without indexes takes longer, and locking the table can prevent other queries from running, affecting application performance.
- **Inefficient joins:** Joins on tables without indexes on the join keys are extremely slow and result in bad query performance.
- **Poor user experience:** Missing indexes can lead to poor user experience in your applications. It can result in slower page loads, application hanging when data is being fetched from the database.

## What is a SARGable query?

SARGable stands for Search Argumentable query, which uses indexes and leads to efficient queries. If a query is SARGable, the database engine quickly locates rows using indexes, avoids scanning the whole table, and improves query performance.

## What are some common mistakes when using GROUP BY?

The common mistakes people encounter when using the **GROUP BY** clause include:

- **Selecting non-aggregated columns not in the GROUP BY clause:** This is a common mistake made by beginners and experts. An example query of this looks like this:

```
sql
```





```
SELECT day, amount FROM Sales
GROUP BY day
```

In the query above, the `amount` column is not part of the `GROUP BY` clause and will throw an error that it must appear in the `GROUP BY` clause. To fix this, you should add an aggregate function to the amount column.

```
sql
```

```
SELECT day, MAX(amount) FROM Sales
GROUP BY day
```

- **Not using aggregate functions:** It is also a common mistake to use `GROUP BY` without aggregate functions. `GROUP BY` usually goes with aggregate functions like `MAX`, `MIN`, `COUNT`, etc.
- **Grouping by multiple columns:** Grouping by multiple columns can make the query meaningless. It is not common to group by many columns, and when this happens, you should check if you really need to group by those columns.

## Why can NOT IN lead to unexpected results with NULLS?

Since NULL is unknown, a `NOT IN` query containing a NULL or NULL in the list of possible values will always return 0 records because of the unknown result introduced by the NULL value. SQL cannot determine for sure whether the value is not in that list.

Let's illustrate this using a table `Sales` that looks like this:

| id | day       | amount |
|----|-----------|--------|
| 1  | Monday    | 200    |
| 2  | Tuesday   | 300    |
| 3  | Wednesday | 600    |
| 4  | Thursday  | 390    |
| 5  | Friday    | 900    |
| 6  | Saturday  | 600    |

If you run the query below, it will return an empty result because SQL cannot determine if the value is not in the list because nothing equals or doesn't equal NULL.

sql



```
SELECT * from sales
WHERE amount NOT IN (200, 300, 600, NULL);
```

## Wrapping up

Mastering SQL queries is essential for anyone working with databases, whether you're a beginner just starting out or an experienced developer looking to sharpen your skills. The 30 questions covered in this guide span from foundational concepts like JOINS and WHERE clauses to advanced topics such as window functions and query optimization.

Remember that SQL proficiency comes with practice. Take time to implement these queries in your own database environment, experiment

with different scenarios, and understand how each function behaves with various data sets. The more you practice, the more confident you'll become in handling complex database challenges during interviews and in real-world applications.

Keep this guide handy as a reference, and don't hesitate to revisit the concepts that challenge you most. Good luck with your SQL journey and upcoming interviews!

### Related Guides



[Is SQL Hard to Learn? \(An Expert's Take\)](#)

[How Long Does It Take to Learn SQL? \(An Expert's Take\)](#)

[SQL vs. Python: Which should you learn for data analysis?](#)

[SQL vs. MySQL: What's the Difference?](#)

---

# Join the Community

roadmap.sh is the [6th most starred project on GitHub](#) and is visited by hundreds of thousands of developers every month.

Rank 6th [out of 28M!](#)

# 351K

GitHub Stars

★ **Star us on GitHub**

Help us reach #1

+90k [every month](#)

# +2.8M

Registered Users

👤 **Register yourself**

Commit to your growth

+2k [every month](#)

# 45K

Discord Members

🗨️ **Join on Discord**

Join the community

[Roadmaps](#) [Guides](#) [FAQs](#) [YouTube](#)

**r.** [roadmap.sh](#) by [@kamrify](#)

Community created roadmaps, best practices, projects, articles, resources and journeys to help you choose your path and grow in your career.

© roadmap.sh · [Terms](#) · [Privacy](#) · [in](#) [📺](#) [X](#) [🦋](#)

## THE NEW STACK

The top DevOps resource for Kubernetes, cloud-native computing, and large-scale development and deployment.

[DevOps](#) · [Kubernetes](#) · [Cloud-Native](#)

[🍪 Cookie Settings](#)